

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import xgboost as xgb
from sklearn.model_selection import train_test_split, StratifiedKFold
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score
import xgboost as xgb
import shap
import warnings
warnings.filterwarnings("ignore")

# Step 1: Load the dataset
def load_data_in_chunks(file_path, nrows=10000):
    """Load a specific number of rows from a large CSV file and add a
    Serial Number column."""
    chunk = pd.read_csv(file_path, nrows=nrows)
    chunk['Serial'] = range(1, len(chunk) + 1) # Adding Serial Number
    column starting from 1
    print(f"Loaded {nrows} rows from {file_path}")
    return chunk

train_data = load_data_in_chunks(r"E:\AYUSH\amex-default-prediction\
train_data.csv", nrows=10000)
test = load_data_in_chunks(r"E:\AYUSH\amex-default-prediction\
test_data.csv", nrows=10000)
train_labels = load_data_in_chunks(r"E:\AYUSH\amex-default-prediction\
train_labels.csv", nrows=10000)
sample_submission = load_data_in_chunks(r"E:\AYUSH\amex-default-
prediction\sample_submission.csv", nrows=10000)

print("\nTop 10 rows of Train Data with Serial Numbers:")
print(train_data.head(10))
print("\nTop 10 rows of Test Data with Serial Numbers:")
print(test.head(10))
print("\nTop 10 rows of Train Data with Serial Numbers:")
print(train_labels.head(10))
print("\nTop 10 rows of Train Data with Serial Numbers:")
print(sample_submission.head(10))

Loaded 10000 rows from E:\AYUSH\amex-default-prediction\train_data.csv
Loaded 10000 rows from E:\AYUSH\amex-default-prediction\test_data.csv
Loaded 10000 rows from E:\AYUSH\amex-default-prediction\
train_labels.csv
Loaded 10000 rows from E:\AYUSH\amex-default-prediction\
sample_submission.csv

```

Top 10 rows of Train Data with Serial Numbers:

	customer_ID	S_2
P_2 \		
0	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-03-09
0.938469		
1	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-04-07
0.936665		
2	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-05-28
0.954180		
3	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-06-13
0.960384		
4	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-07-16
0.947248		
5	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-08-04
0.945964		
6	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-09-18
0.940705		
7	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-10-08
0.914767		
8	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-11-20
0.950845		
9	0000099d6bd597052cdcda90ffabf56573fe9d7c79be5f...	2017-12-04
0.868580		

	D_39	B_1	B_2	R_1	S_3	D_41
B_3 ... \						
0	0.001733	0.008724	1.006838	0.009228	0.124035	0.008771
0.004709 ...						
1	0.005775	0.004923	1.000653	0.006151	0.126750	0.000798
0.002714 ...						
2	0.091505	0.021655	1.009672	0.006815	0.123977	0.007598
0.009423 ...						
3	0.002455	0.013683	1.002700	0.001373	0.117169	0.000685
0.005531 ...						
4	0.002483	0.015193	1.000727	0.007605	0.117325	0.004653
0.009312 ...						
5	0.001746	0.007863	1.005006	0.004220	0.110946	0.009857
0.009866 ...						
6	0.002183	0.018859	1.008024	0.004509	0.103329	0.006603
0.000783 ...						
7	0.003029	0.014324	1.000242	0.000263	0.108115	0.009527
0.007836 ...						
8	0.009896	0.016888	1.003995	0.001789	0.102792	0.002519
0.009817 ...						
9	0.001082	0.001930	1.007504	0.001772	0.100470	0.004626
0.006073 ...						

	D_137	D_138	D_139	D_140	D_141	D_142	D_143
D_144 \							

0	NaN	NaN	0.002427	0.003706	0.003818	NaN	0.000569
0.000610							
1	NaN	NaN	0.003954	0.003167	0.005032	NaN	0.009576
0.005492							
2	NaN	NaN	0.003269	0.007329	0.000427	NaN	0.003429
0.006986							
3	NaN	NaN	0.006117	0.004516	0.003200	NaN	0.008419
0.006527							
4	NaN	NaN	0.003671	0.004946	0.008889	NaN	0.001670
0.008126							
5	NaN	NaN	0.001924	0.008598	0.004529	NaN	0.000674
0.002223							
6	NaN	NaN	0.001336	0.004361	0.009387	NaN	0.007727
0.007661							
7	NaN	NaN	0.002397	0.008452	0.005553	NaN	0.001831
0.009616							
8	NaN	NaN	0.009742	0.003968	0.007945	NaN	0.008722
0.004369							
9	NaN	NaN	0.003611	0.009607	0.007266	NaN	0.008763
0.004753							

	D_145	Serial
0	0.002674	1
1	0.009217	2
2	0.002603	3
3	0.009600	4
4	0.009827	5
5	0.002884	6
6	0.002225	7
7	0.007385	8
8	0.000995	9
9	0.009068	10

[10 rows x 191 columns]

Top 10 rows of Test Data with Serial Numbers:

	customer_ID	S_2
P_2 \		
0	00000469ba478561f23a92a868bd366de6f6527a684c9a...	2019-02-19
0.631315		
1	00000469ba478561f23a92a868bd366de6f6527a684c9a...	2019-03-25
0.587042		
2	00000469ba478561f23a92a868bd366de6f6527a684c9a...	2019-04-25
0.609056		
3	00000469ba478561f23a92a868bd366de6f6527a684c9a...	2019-05-20
0.614911		
4	00000469ba478561f23a92a868bd366de6f6527a684c9a...	2019-06-15
0.591673		
5	00000469ba478561f23a92a868bd366de6f6527a684c9a...	2019-07-13
0.587472		

```

6 00000469ba478561f23a92a868bd366de6f6527a684c9a... 2019-08-16
0.625006
7 00000469ba478561f23a92a868bd366de6f6527a684c9a... 2019-09-29
0.597074
8 00000469ba478561f23a92a868bd366de6f6527a684c9a... 2019-10-12
0.568930
9 00001bf2e77ff879fab36aa4fac689b9ba411dae63ae39... 2018-04-22
0.894195

```

	D_39	B_1	B_2	R_1	S_3	D_41
B_3 ... \						
0 0.001912	0.010728	0.814497	0.007547	0.168651	0.009971	
0.002347 ...						
1 0.005275	0.011026	0.810848	0.001817	0.241389	0.000166	
0.009132 ...						
2 0.003326	0.016390	1.004620	0.000114	0.266976	0.004196	
0.004192 ...						
3 0.009065	0.021672	0.816549	0.009722	0.188947	0.004123	
0.015325 ...						
4 0.238794	0.015923	0.810456	0.002026	0.180035	0.000731	
0.011281 ...						
5 0.005827	0.007959	1.001009	0.008814	0.173701	0.000653	
0.009779 ...						
6 0.238704	0.013420	0.816605	0.009849	0.170644	0.002495	
0.019999 ...						
7 0.003537	0.017426	1.000670	0.005143	0.158006	0.000985	
0.011962 ...						
8 0.121385	0.010779	1.009347	0.006923	0.149413	0.000396	
0.003576 ...						
9 0.325657	0.020970	1.001803	0.005125	0.073243	0.005347	
0.001597 ...						

	D_137	D_138	D_139	D_140	D_141	D_142	D_143
D_144 \							
0 NaN	NaN	NaN	0.004669	NaN	NaN	NaN	
0.008281							
1 NaN	NaN	0.000142	0.004940	0.009021	NaN	0.003695	
0.003753							
2 NaN	NaN	0.000074	0.002114	0.004656	NaN	0.003155	
0.002156							
3 NaN	NaN	0.004743	0.006392	0.002890	NaN	0.006044	
0.005206							
4 NaN	NaN	0.008133	0.004329	0.008384	NaN	0.001008	
0.007421							
5 NaN	NaN	0.007001	0.003962	0.005530	NaN	0.009870	
0.009667							
6 NaN	NaN	0.001132	0.007676	0.005410	NaN	0.006762	
0.005664							
7 NaN	NaN	0.008730	0.000628	0.002821	NaN	0.004141	

```

0.005733
8      NaN      NaN  0.005912  0.001250  0.006543      NaN  0.009160
0.003690
9      NaN      NaN  0.008065  0.009861  0.009535      NaN  0.003964
0.008436

```

```

      D_145  Serial
0         NaN      1
1  0.001460      2
2  0.006482      3
3  0.007855      4
4  0.009471      5
5  0.005398      6
6  0.002627      7
7  0.005657      8
8  0.003219      9
9  0.008323     10

```

[10 rows x 191 columns]

Top 10 rows of Train Data with Serial Numbers:

	customer_ID	target	Serial
0	0000099d6bd597052cdcd90ffabf56573fe9d7c79be5f...	0	1
1	00000fd6641609c6ece5454664794f0340ad84dddce9a2...	0	2
2	00001b22f846c82c51f6e3958ccd81970162bae8b007e8...	0	3
3	000041bdba6ecadd89a52d11886e8eaaec9325906c9723...	0	4
4	00007889e4fcd2614b6cbe7f8f3d2e5c728eca32d9eb8a...	0	5
5	000084e5023181993c2e1b665ac88dbb1ce9ef621ec537...	0	6
6	000098081fde4fd64bc4d503a5d6f86a0aedc425c96f52...	0	7
7	0000d17a1447b25a01e42e1ac56b091bb7cbb06317be4c...	0	8
8	0000f99513770170a1aba690daeeb8a96da4a39f11fc27...	1	9
9	00013181a0c5fc8f1ea38cd2b90fe8ad2fa8cad9d9f13e...	1	10

Top 10 rows of Train Data with Serial Numbers:

	customer_ID	prediction
Serial		
0	00000469ba478561f23a92a868bd366de6f6527a684c9a...	0
1		
1	00001bf2e77ff879fab36aa4fac689b9ba411dae63ae39...	0
2		
2	0000210045da4f81e5f122c6bde5c2a617d03eef67f82c...	0
3		
3	00003b41e58ede33b8daf61ab56d9952f17c9ad1c3976c...	0
4		
4	00004b22eaeecb0ec976890c1d9bfc14fd9427e98c4ee9...	0
5		
5	00004ffe6e01e1b688170bbd108da8351bc4c316eacfef...	0
6		
6	00007cfcce97abfa0b4fa0647986157281d01d3ab90de9...	0
7		

```

7  000089cc2a30dad8e6ba39126f9d86df6088c9f975093a...      0
8
8  00008f50a1dd76fa211ba36a2b0d5a1b201e4134a5fd53...      0
9
9  0000b48a4f27dc1d61e78d081678e811620300b88eb3ab...      0
10

# Step 2: Exploratory Data Analysis (EDA)
def explore_data(data):
    """Perform basic exploration to understand the data."""
    print("Basic Information of the Dataset:")
    print("-" * 50)
    print(data.info())
    print("\nBasic Statistics of Numerical Columns:")
    print("-" * 50)
    print(data.describe().T)

    print("\nMissing Values in Top Columns:")
    print("-" * 50)
    print(data.isnull().sum().sort_values(ascending=False).head(10))

    # Check if 'target' column exists before analyzing its
    # distribution
    if 'target' in data.columns:
        print("\nTarget Distribution:")
        print("-" * 50)
        print(data['target'].value_counts(normalize=True))

        # Visualizing target distribution
        sns.countplot(x='target', data=data)
        plt.title("Target Distribution")
        plt.xlabel("Target Class")
        plt.ylabel("Count")
        plt.show()
    else:
        print("\n'Target' column is not present in the dataset.")

# Call the function
explore_data(train_data)

Basic Information of the Dataset:
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Columns: 191 entries, customer_ID to Serial
dtypes: float64(185), int64(2), object(4)
memory usage: 14.6+ MB
None

Basic Statistics of Numerical Columns:

```

-----					
	count	mean	std	min	
25% \					
P_2	9936.0	0.650498	0.252416	-2.569212e-01	0.471264
D_39	10000.0	0.157031	0.275827	8.701630e-07	0.004574
B_1	10000.0	0.126147	0.212428	-1.414690e-01	0.009126
B_2	10000.0	0.617122	0.403145	3.432000e-05	0.091656
R_1	10000.0	0.074035	0.219977	3.031180e-06	0.002863
...	...	...	...	...	...
D_142	1572.0	0.370263	0.249771	-8.804185e-03	0.150429
D_143	9847.0	0.162935	0.364742	2.172557e-06	0.002939
D_144	9937.0	0.046364	0.174197	7.773110e-07	0.002674
D_145	9847.0	0.057846	0.226945	1.577788e-07	0.002966
Serial	10000.0	5000.500000	2886.895680	1.000000e+00	2500.750000

	50%	75%	max
P_2	0.690195	0.866060	1.009926
D_39	0.009296	0.238015	4.268383
B_1	0.032948	0.122675	1.323411
B_2	0.814048	1.002262	1.009999
R_1	0.005726	0.008528	2.259283
...	...	...	...
D_142	0.352169	0.559880	1.185992
D_143	0.005944	0.008907	1.010000
D_144	0.005323	0.008105	1.342362
D_145	0.005932	0.008878	4.282032
Serial	5000.500000	7500.250000	10000.000000

[187 rows x 8 columns]

Missing Values in Top Columns:

-----	
D_87	9990
D_88	9980
D_111	9975
D_110	9975
B_39	9972
D_108	9944
D_73	9926
B_42	9831

```
D_138      9628
D_136      9628
dtype: int64
```

'Target' column is not present in the dataset.

*# Step 3: Handling Missing Values*

```
def handle_missing_values(train, test):
    """Fill missing values with a placeholder."""
    train.fillna(-999, inplace=True)
    test.fillna(-999, inplace=True)
    return train, test
```

```
train, test = handle_missing_values(train_data, test)
```

*# Step 4: Encoding Categorical Variables*

```
from sklearn.preprocessing import LabelEncoder
```

```
def encode_categorical(train, test):
    """Label encode categorical features."""
    categorical_cols = train.select_dtypes(include=['object']).columns
    encoder = LabelEncoder()
```

*# Apply encoding for each categorical column*

```
for col in categorical_cols:
```

```
    # Fit on training data and transform both train and test sets
    train[col] = encoder.fit_transform(train[col].astype(str))
```

*# Transform test data using the same encoder*

*# We use 'fit\_transform' only on training data to avoid test data leakage*

```
    test[col] = encoder.transform(test[col].astype(str))
```

```
    return train, test
```

*# Apply the encoding function to train and test datasets*

```
train, test = encode_categorical(train_data, test)
```

```
from sklearn.preprocessing import StandardScaler
```

```
def scale_features(train, test):
    """Standardize numerical features with aligned columns."""
    scaler = StandardScaler()
```

*# Drop non-numeric columns first*

```
train_numeric = train.select_dtypes(include=['int64', 'float64'])
test_numeric = test.select_dtypes(include=['int64', 'float64'])
```

*# Align columns between train and test*

```
train_numeric, test_numeric = train_numeric.align(test_numeric,
join='inner', axis=1)
```



```

# Fit scaler on train data and transform both train and test
train_scaled = scaler.fit_transform(train_numeric)
test_scaled = scaler.transform(test_numeric)

# Replace the numerical columns in the original train and test
DataFrames
train[train_numeric.columns] = train_scaled
test[train_numeric.columns] = test_scaled

return train, test

# Apply feature scaling
train, test = scale_features(train_data, test)

print("Feature scaling completed successfully!")

```

Feature scaling completed successfully!

*# Step 6: Splitting Data*

```
from sklearn.model_selection import train_test_split
```

```

def split_train_data(train):
    """Split data into training and validation sets."""
    if not isinstance(train, pd.DataFrame):
        raise ValueError("Input should be a DataFrame.")

    X = train.drop(columns=['target', 'id'], errors='ignore')
    y = train['target'] if 'target' in train.columns else None

    if y is None:
        raise ValueError("Column 'target' not found in the
DataFrame.")

    X_train, X_val, y_train, y_val = train_test_split(X, y,
test_size=0.2, stratify=y, random_state=42)
    return X_train, X_val, y_train, y_val

```

*# Step 7: Custom Evaluation Metric*

```

def amex_metric(y_true, y_pred):
    """Define the AMEX evaluation metric."""
    def top_four_percent_captured(y_true, y_pred):
        df = pd.DataFrame({'y_true': y_true, 'y_pred': y_pred})
        df = df.sort_values('y_pred', ascending=False)
        df['weight'] = df['y_true'].apply(lambda x: 20 if x == 1 else
1)

        cutoff = int(0.04 * df['weight'].sum())
        return df.iloc[:cutoff]['y_true'].sum() / df['y_true'].sum()

```

```

    gini = 2 * roc_auc_score(y_true, y_pred) - 1
    return 0.5 * (gini + top_four_percent_captured(y_true, y_pred))

# Step 8: Train Baseline Logistic Regression
def split_train_data(train):
    """Split data into training and validation sets."""
    # Check for available columns
    print("Columns in train:", train.columns)

    # Drop 'target' or 'id' only if they exist
    drop_cols = [col for col in ['target', 'id'] if col in
train.columns]
    X = train.drop(columns=drop_cols) # Drop valid columns only

    y = train['target'] if 'target' in train.columns else None #
Ensure 'target' exists

    if y is None:
        raise ValueError("Column 'target' not found in the dataset!")

    X_train, X_val, y_train, y_val = train_test_split(
        X, y, test_size=0.2, stratify=y, random_state=42
    )
    return X_train, X_val, y_train, y_val

# Step 9: Train Advanced XGBoost Model
from sklearn.model_selection import train_test_split

def split_train_data(train):
    target_column = 'target' # Update this to the correct column name
if needed
    if target_column not in train.columns:
        raise ValueError(f"Column '{target_column}' not found in the
dataset!")

    y = train[target_column]
    X = train.drop(columns=[target_column])

    X_train, X_val, y_train, y_val = train_test_split(
        X, y, test_size=0.2, stratify=y, random_state=42
    )
    return X_train, X_val, y_train, y_val

# Step 10: Analyze Feature Importance
# Ensure 'customer_ID' has the same data type in both datasets
train_data["customer_ID"] = train_data["customer_ID"].astype(str)

```

```

train_labels["customer_ID"] = train_labels["customer_ID"].astype(str)

# Check if 'customer_ID' is unique before setting it as the index
if train_data["customer_ID"].is_unique and
train_labels["customer_ID"].is_unique:
    # Set 'customer_ID' as the index for both DataFrames before
    concatenation
    train = pd.concat([train_data.set_index("customer_ID"),
                        train_labels.set_index("customer_ID")],
axis=1).reset_index()
else:
    # If 'customer_ID' is not unique, merge based on 'customer_ID'
    instead
    train = pd.merge(train_data, train_labels, on="customer_ID",
how="inner")

# Prepare data for training
X = train.drop(columns=["customer_ID", "target"], errors="ignore") #
Features
y = train["target"] # Target column

# Ensure that X and y are not empty
if len(X) > 0 and len(y) > 0:
    # Train-validation split with a valid test_size and train_size
    X_train, X_val, y_train, y_val = train_test_split(X, y,
test_size=0.2, random_state=42)

    # Step 5: Train XGBoost Model
    xgb_model = xgb.XGBClassifier(
        objective="binary:logistic",
        eval_metric="auc",
        learning_rate=0.05,
        max_depth=6,
        subsample=0.8,
        colsample_bytree=0.8,
        random_state=42
    )
    xgb_model.fit(X_train, y_train, eval_set=[(X_val, y_val)],
early_stopping_rounds=50, verbose=True)

# Step 6: Feature Importance Visualization
def analyze_features(model, X_val):
    """
    Visualize feature importance using SHAP and XGBoost.
    """
    # XGBoost Feature Importance
    xgb.plot_importance(model)
    plt.title("XGBoost Feature Importance")
    plt.show()

```

```

    # SHAP Feature Importance
    explainer = shap.Explainer(model, X_val)
    shap_values = explainer(X_val)
    shap.summary_plot(shap_values, X_val)

    # Step 7: Call the function
    analyze_features(xgb_model, X_val)

else:
    print("Error: The input data is empty. Please check the data.")

Error: The input data is empty. Please check the data.

# Step 11: Generate Submission File
def encode_categorical(train, test):
    """Label encode categorical features."""
    categorical_cols = train.select_dtypes(include=['object']).columns
    encoder = LabelEncoder()
    for col in categorical_cols:
        train[col] = encoder.fit_transform(train[col].astype(str))
        test[col] = encoder.transform(test[col].astype(str))
    return train, test

# Step 2: Train XGBoost Model without early stopping
def train_xgb_model(train_data, train_labels):
    """Train XGBoost model on the training data."""
    # Prepare training data
    X = train_data.drop(columns=['target'], errors='ignore') # Drop
'target' if it exists
    y = train_labels['target']

    # Split into train and validation sets
    X_train, X_val, y_train, y_val = train_test_split(X, y,
test_size=0.2, random_state=42)

    # Initialize the XGBoost classifier
    xgb_model = xgb.XGBClassifier(
        objective="binary:logistic",
        learning_rate=0.05,
        max_depth=6,
        subsample=0.8,
        colsample_bytree=0.8,
        random_state=42
    )

    # Train the model without early stopping
    xgb_model.fit(X_train, y_train, verbose=True)

    return xgb_model

```

```

# Step 3: Generate Submission File
def generate_submission(test, xgb_model):
    """Create the final submission file."""
    # Check if the 'id' column exists in the test DataFrame
    if 'id' not in test.columns:
        print("Error: 'id' column not found in the test DataFrame.")
        return

    # Prepare test data (drop 'id' column and other non-predictive
    columns)
    X_test = test.drop(columns=['id'], errors='ignore')

    # Convert test data to DMatrix format for XGBoost
    dtest = xgb.DMatrix(X_test)

    # Make predictions (ensure binary classification, or adjust as
    necessary)
    predictions = xgb_model.predict(dtest)

    # If the model is binary classification, convert probabilities to
    0 or 1
    if predictions.shape[0] > 1: # In case of output being
    probabilities
        predictions = (predictions > 0.5).astype(int) # Threshold at
    0.5

    # Create the submission DataFrame
    submission = pd.DataFrame({
        'id': test['id'],
        'prediction': predictions
    })

    # Save the submission file to the specified path
    submission.to_csv(r'E:\AYUSH\amex-default-prediction\
    sample_submission.csv', index=False)
    print("Submission file saved as sample_submission.csv")

# Example usage:
# Assuming 'train_data', 'train_labels', and 'test' are already loaded
# with the respective data

# Step 4: Encode categorical features (if needed)
train_data, test = encode_categorical(train_data, test)

# Step 5: Train the XGBoost model
xgb_model = train_xgb_model(train_data, train_labels) # Ensure
'train_data' and 'train_labels' are defined

# Step 6: Generate the submission file

```

```
generate_submission(test, xgb_model)
```

```
Error: 'id' column not found in the test DataFrame.
```